

# GNU Assembler



CMPE230 - Spring'24

Gökçe Uludoğan

# GNU Assembler

- the assembler developed by the GNU Project.
- used to assemble the GNU operating system and the Linux kernel.
- uses AT&T assembly syntax.

# AT&T Syntax

- Similar to any other assembler syntax.
- Consists of a series of directives, labels, instructions
- Composed of a mnemonic followed by a maximum of three operands
  - **the ordering of the operands are reversed.**

Intel Syntax (e.g. A86)	mnemonic	destination, source
AT&T Syntax	mnemonic	source, destination

# Registers

Register Names

64-bit register	32-bit sub-register	16-bit sub-register	8-bit sub-register
%rax	%eax	%ax	%al
%rbx	%ebx	%bx	%bl
%rcx	%ecx	%cx	%cl
%rdx	%edx	%dx	%dl
%rsi	%esi	%si	%sil
%rdi	%edi	%di	%dil
%rbp	%ebp	%bp	%bpl
%rsp	%esp	%sp	%spl
%r8	%r8d	%r8w	%r8b
%r9	%r9d	%r9w	%r9b
%r10	%r10d	%r10w	%r10b
%r11	%r11d	%r11w	%r11b
%r12	%r12d	%r12w	%r12b
%r13	%r13d	%r13w	%r13b
%r14	%r14d	%r14w	%r14b
%r15	%r15d	%r15w	%r15b

# Registers

## Registers

<b>%rip</b>	Instruction pointer
<b>%rsp</b>	Stack pointer
<b>%rax</b>	Return value
<b>%rdi</b>	1st argument
<b>%rsi</b>	2nd argument
<b>%rdx</b>	3rd argument
<b>%rcx</b>	4th argument
<b>%r8</b>	5th argument
<b>%r9</b>	6th argument
<b>%r10,%r11</b>	Callee-owned
<b>%rbx,%rbp, %r12-%15</b>	Caller-owned

## Instruction suffixes

<b>b</b>	byte
<b>w</b>	word (2 bytes)
<b>l</b>	long /doubleword (4 bytes)
<b>q</b>	quadword (8 bytes)

Suffix is elided when can be inferred from operands. e.g. operand **%rax** implies **q**, **%eax** implies **l**, and so on

Mnemonic	Purpose	Examples
<code>mov <i>src,dest</i></code>	Move data between registers, load immediate data into registers, move data between registers and memory.	<code>mov \$4,%eax # Load constant into eax</code> <code>mov %eax,%ebx # Copy eax into ebx</code> <code>mov %ebx,123 # Copy ebx to mem. 123</code>
<code>push <i>src</i></code>	Insert a value onto the stack. Useful for passing arguments, saving registers, etc.	<code>push %ebp</code>
<code>pop <i>dest</i></code>	Remove topmost value from the stack. Equivalent to " <code>mov (%esp),<i>dest</i></code> ; add \$4,%esp"	<code>pop %ebp</code>
<code>call <i>func</i></code>	Push the address of the next instruction and start executing <i>func</i> .	<code>call print_int</code>
<code>ret</code>	Pop the return program counter, and jump there. Ends a subroutine.	<code>ret</code>
<code>add <i>src,dest</i></code>	<code><i>dest=dest+src</i></code>	<code>add %ebx, %eax # Add ebx to eax</code>
<code>mul <i>src</i></code>	Multiply <i>eax</i> and <i>src</i>	<code>mul %ebx #Multiply eax by ebx</code>
<code>jmp <i>label</i></code> <code>j<sub>l</sub> <i>label</i></code>	Goto the instruction <i>label</i> :. Skips anything else in the way. Goto <i>label</i> if the comparison came out as less-than. Others: <code>jle (&lt;=)</code> , <code>je (==)</code> , <code>jge (&gt;=)</code> , <code>jg (&gt;)</code> , <code>jne (!=)</code> , and so on.	<code>jmp post_mem</code> <code>mov %eax,0 # Write to NULL!</code> <code>post_mem: # OK here...</code>
<code>cmp <i>a,b</i></code>	Compare two values. Sets flags that are used by the conditional jumps (below).	<code>cmp \$10,%eax</code>

# GDB

- The GNU Debugger
- Install via

```
apt-get update  
apt-get install gdb
```

- Run with

```
gcc <file>.s -c -g  
ld <file>.o -o <program>  
gdb <program>
```

# GDB

`r` for running the program

```
(gdb) r
Starting program: /root/ps9/ex3

12345678
[Inferior 1 (process 21656) exited normally]
```



# GDB

`b` for setting a breakpoint

`s` for step

```
Reading symbols from ex3...
(gdb) br 1
Breakpoint 1 at 0x401000: file ex3.s, line 13.
(gdb) r
Starting program: /root/ps9/ex3

Breakpoint 1, _start () at ex3.s:13
13      lea input_buffer(%rip), %rsi
(gdb) s
14      mov $256, %edx
(gdb) s
15      call input_fn
```

# GDB TUI

[https://ftp.gnu.org/old-gnu/Manuals/gdb/html\\_chapter/gdb\\_19.html](https://ftp.gnu.org/old-gnu/Manuals/gdb/html_chapter/gdb_19.html)

> gdb -tui <program>

```
ex3.s
21     call print_fn
22
23     call exit_fn
24
25 input_fn:
26     mov $0, %rax           # System call number for sys_read (0)
27     mov $0, %rdi          # File descriptor for standard input (0)
> 28     syscall
29     ret
30
31 print_fn:
32     mov $1, %rax          # System call number for sys_write (1)

native process 26553 (src) In: input_fn                                L28    PC: 0x40103a
Starting program: /root/ps9/ex3

Breakpoint 1, _start () at ex3.s:13
(gdb) s
input_fn () at ex3.s:26
(gdb) s
(gdb) |
```

# GDB TUI Layouts

tui layout regs

tui layout asm

[https://ftp.gnu.org/old-gnu/Manuals/gdb/html\\_chapter/gdb\\_19.html](https://ftp.gnu.org/old-gnu/Manuals/gdb/html_chapter/gdb_19.html)

Register group: general

rax	0x0	0	rbx	0x0	0
rcx	0x0	0	rdx	0x100	256
rsi	0x40200a	4202506	rdi	0x0	0
rbp	0x0	0x0	rsp	0x7fffffff128	0x7fffffff128
r8	0x0	0	r9	0x0	0

```
> 27     mov $0, %rdi           # File descriptor for standard input (0)
    28     syscall
    29     ret
    30
    31 print_fn:
    32     mov $1, %rax           # System call number for sys_write (1)
```

native process 26553 (regs) In: input\_fn

L28 PC: 0x40103a

input\_fn () at ex3.s:26

(gdb) s

(gdb) assembly

Undefined command: "assembly". Try "help".

(gdb) layout next

(gdb) layout next

(gdb)

# Example

- System calls

[https://blog.rchapman.org/posts/Linux\\_System\\_Call\\_Table\\_for\\_x86\\_64/](https://blog.rchapman.org/posts/Linux_System_Call_Table_for_x86_64/)

```
.section .data
hello: .string "12345678\n"

.section .bss
input_buffer: .space 256      # Reserve 256 bytes for input buffer

.section .text
.global _start

_start:

    mov $0, %rax              # System call number for sys_read (0)
    mov $0, %rdi              # File descriptor for standard input (0)
    lea input_buffer(%rip), %rsi
    mov $256, %edx
    syscall

    mov $1, %rax              # System call number for sys_write (1)
    mov $1, %rdi              # File descriptor for standard output (1)
    syscall

    lea hello(%rip), %rsi
    mov $9, %edx

    mov $1, %rax              # System call number for sys_write (1)
    mov $1, %rdi              # File descriptor for standard output (1)
    syscall

    mov $60, %rax
    mov $0, %rdi
    syscall
```

# Example

```
.section .data
hello: .string "12345678\n"

.section .bss
input_buffer: .space 256      # Reserve 256 bytes for input buffer

.section .text
.global _start

_start:
    lea input_buffer(%rip), %rsi
    mov $256, %edx
    call input_fn

    call print_fn

    lea hello(%rip), %rsi
    mov $9, %edx
    call print_fn

    call exit_fn

input_fn:
    mov $0, %rax                # System call number for sys_read (0)
    mov $0, %rdi                # File descriptor for standard input (0)
    syscall
    ret

print_fn:
    mov $1, %rax                # System call number for sys_write (1)
    mov $1, %rdi                # File descriptor for standard output (1)
    syscall
    ret

exit_fn:
    mov $60, %rax
    mov $0, %rdi
    syscall
```

# Reverse string

```
.section .bss
input_buffer: .space 10      # Allocate 256 bytes for input buffer
output_buffer: .space 10

.section .text
.global _start

_start:
    # Read input from standard input
    mov $0, %eax             # syscall number for sys_read
    mov $0, %edi             # file descriptor 0 (stdin)
    lea input_buffer(%rip), %rsi # pointer to the input buffer
    mov $10, %edx            # maximum number of bytes to read
    syscall                  # perform the syscall

    lea input_buffer(%rip), %r15
    lea output_buffer(%rip), %r14
    add $2, %r14
    mov $0, %r13

reverse:
    mov $0, %rbx
    movb (%r15), %bl
    cmp $'\n', %rbx          # Check if we've parsed all input
    je print_output_buffer    # If zero, finish the loop
    inc %r13

    mov %bl, (%r14)
    inc %r15
    dec %r14
    jmp reverse
```

```
print_output_buffer:
    lea output_buffer(%rip), %rsi
    mov %r13, %rdx
    mov $1, %eax             # syscall number for sys_write
    mov $1, %edi             # file descriptor 1 (stdout)
    syscall

exit_program:
    # Exit the program
    mov $60, %eax            # syscall number for sys_exit
    xor %edi, %edi           # exit code 0
    syscall
```

# Reverse string using stack

```
.section .bss
input_buffer: .space 10      # Allocate 256 bytes for input buffer
output_buffer: .space 10

.section .text
.global _start

_start:

    # Read input from standard input
    mov $0, %eax             # syscall number for sys_read
    mov $0, %edi             # file descriptor 0 (stdin)
    lea input_buffer(%rip), %rsi # pointer to the input buffer
    mov $10, %edx            # maximum number of bytes to read
    syscall                  # perform the syscall

    lea input_buffer(%rip), %r15
    lea output_buffer(%rip), %r14
    mov $0, %r13

reverse:
    mov $0, %rbx
    movb (%r15), %bl
    cmp $'\n', %rbx
    je store                 # Check if we've parsed all input
                                # If zero, finish the loop
    push %rbx
    inc %r13
    inc %r15
    jmp reverse
```

```
store:
    cmp $0, %r13
    je print_output_buffer
    dec %r13
    pop %rbx
    movb %bl, (%r14)
    inc %r14
    jmp store

print_output_buffer:
    lea output_buffer(%rip), %rsi
    mov %r14, %rdx
    mov $1, %eax             # syscall number for sys_write
    mov $1, %edi             # file descriptor 1 (stdout)
    syscall

exit_program:
    # Exit the program
    mov $60, %eax            # syscall number for sys_exit
    xor %edi, %edi           # exit code 0
    syscall
```

# Output?

```

    .section .data

n: .long 50

.section .text
.global _start

_start:
    mov $n, %rsi
    mov $1, %rdx

    call print_fn
    call exit_fn

print_fn:
    mov $1, %rax           # System call number for sys_write (1)
    mov $1, %rdi           # File descriptor for standard output (1)
    syscall
    ret

exit_fn:
    mov $60, %rax
    mov $0, %rdi
    syscall
```